

ECE448 Spring 2022

Assignment 1: Search

In this assignment, you will build general-purpose search algorithms and apply them to solving puzzles. In Part 1, you will be in charge of a "Pacman"-like agent that needs to find a path through maze to eat a dot or "food pellet." In Part 2, you will need to find a single path that goes through all the dots in the maze.




Image courtesy of the ECE448 course lecturer (Prof. Hasegawa-Johnson) at UIUC

Programming language

This MP will be written in Python. If you've never used Python before, you should start getting used to it. A good place to start is the [Python Tutorial](#) (also available in [hardcopy form](#)). You should install version 3.6 or 3.7 on your computer, as well as the [pygame](#) graphics package.

Your code may import extra modules, but only ones that are part of the [standard python library](#). Unless otherwise specified in the instructions for a specific MP, the only external library available during our grading process will be pygame. For example: in mp1, numpy is not allowed.

Part 1: Basic Pathfinding

Consider the problem of finding the shortest path from a given start state while eating one or more dots or "food pellets." The image at the top of this page illustrates the simple scenario of a single dot, which in this case can be viewed as the unique goal state. The maze layout will be given to you in a simple text format, where '%' stands for walls, 'P' for the starting position, and '.' for the dot(s) (see sample [maze.txt](#)). All step costs are equal to one.

Implement the state representation, transition model, and goal test needed for solving the problem in the general case of multiple dots. For the state representation, besides your current position in the maze, is there anything else you need to keep track of? For the goal test, keep in mind that in the case of multiple dots, the Pacman does not necessarily have a unique ending position. Next, implement a unified top-level search routine that can work with all of the following search strategies, as covered in class and/or the textbook (Note: you don't need to find the best solution when using DFS):

- Depth-first search
- Breadth-first search
- Greedy best-first search
- A* search

For this part of the assignment, use the Manhattan distance from the current position to the goal as the

heuristic function for greedy and A* search.

Run each of the four search strategies on the following inputs:

- Medium maze (mediumMaze.txt)
- Big maze (bigMaze.txt)
- Open maze (openMaze.txt)

The provided code will generate a pretty picture of your solution. Your report should include

- The solution pictures.
- The length of the path. Include both the start and goal positions as part of your path and path length.
- Number of nodes expanded by the search algorithm.

Part 2: Search with multiple dots

Now consider the harder problem of finding the shortest path through a maze while hitting multiple dots. Once again, the Pacman is initially at P, but now there is no single goal position. Instead, the goal is achieved whenever the Pacman manages to eat all the dots. Once again, we assume unit step costs.

As instructed in Part 1, your state representation, goal test, and transition model should already be adapted to deal with this scenario. The next challenge is to solve the following inputs using A* search using an admissible heuristic designed by you:

- Tiny search (tinySearch.txt)
- Small search (smallSearch.txt)
- Medium search (mediumSearch.txt)

You should be able to handle the tiny search using uninformed BFS. In fact, it is a good idea to try that first for debugging purposes, to make sure your representation works with multiple dots. However, to successfully handle all the inputs, it is crucial to come up with a good heuristic. For full credit, your heuristic should be admissible and should permit you to find the solution for the medium search in a reasonable amount of time. In your report, explain the heuristic you chose, and discuss why it is admissible and whether it leads to an optimal solution.

Hints: In the past almost all working solutions to this problem have used a heuristic based on the minimum spanning tree. The minimum spanning tree of a set of points can be computed easily via Kruskal's algorithm or Prim's algorithm. If T is the total length of the edges in the minimum spanning tree, then the shortest path connecting all the points must have length between T and $2T$. Now, suppose you are at some location (x,y) with a set of S dots still to reach. Your heuristic function h might be the sum of the distance from (x,y) to the nearest dot, plus the MST length for the dots in S .

Extra Credit Suggestion

Sometimes, even with A* and a good heuristic, finding the optimal path through all the dots is hard. In these cases, we'd still like to find a reasonably good path, quickly. Write a suboptimal search algorithm that will do a good job on this big maze (bigDots.txt). Your algorithm could either be A* with a non-admissible heuristic, or something different altogether. In your report, discuss your approach and output the solution cost and number of expanded nodes. Note that the extra credit will be capped to 10% of what the assignment is worth.

Provided Code Skeleton

We have provided (.tar file or .zip file) all the code to get you started on your MP, which means you will only have to write the search functions. Do not modify provided code. You will only have to modify search.py.

maze.py

- `getStart()` :- Returns a tuple of the starting position, (row, col)
- `getObjectives()` :- Returns a list of tuples that correspond to the dot positions, [(row1, col1), (row2, col2)]
- `isValidMove(row, col)` :- Returns the boolean True if the (row, col) position is valid. Returns False otherwise.
- `getNeighbors(row, col)` :- Given a position, returns the list of tuples that correspond to valid neighbor positions. This will return at most 4 neighbors but may return less.

There are 4 methods to implement in this file, namely `bfs(maze)`, `dfs(maze)`, `greedy(maze)`, and `astar(maze)`. (You may need to add another named search method if you implement an additional search method for extra credit.) Each of these functions takes in a maze instance and should return both the path taken (as a list of tuples) and the number of states explored. The maze instance provided will already be instantiated, and the above methods will be accessible.

To understand how to run the MP, read the provided README.md or run `python3 mp1.py -h` into your terminal. **The following command will display a maze and let you create a path manually using the arrow keys.**

```
python mp1.py --human maps/maze.txt
```

The following command will run your astar search method on the maze.

```
python mp1.py --method astar maps/maze.txt
```

You can also save your output picture as a file in tga format. If your favorite document formatter doesn't handle tga, tools such as gimp can convert it to other formats (e.g. jpg).

Tips

- In your implementation, make sure you get all the bookkeeping right. This includes handling of repeated states (in particular, what happens when you find a better path to a state already on the frontier) and saving the optimal solution path.
- Pay attention to tie breaking. If you have multiple nodes on the frontier with the same minimum value of the evaluation function, the speed of your search and the quality of the solution may depend on which one you select for expansion.
- Implement all four strategies using a similar approach and coding style. In particular, while DFS can be implemented very compactly using recursion, you must store the frontier in an explicit stack, queue or priority queue (depending on the search algorithm) for this assignment. Among other things, limits on recursion depth can be (depending on your installation) much lower than the number of objects that you can pack into an explicit stack.
- You will be graded primarily on the correctness of your solution, not on the efficiency and elegance of your data structures. For example, we don't care whether your priority queue or repeated state detection uses brute-force search, as long as you end up expanding exactly the correct number of

nodes (except for small differences caused by differences among tie-breaking strategies) and find the optimal solution. So, feel free to use "dumb" data structures as long as it makes your life easier and still enables you to find the solutions to all the inputs in a reasonable amount of time.

Deliverables

This MP will be submitted via blackboard.

Please upload **only the following files** to blackboard.

1. search.py - your solution python file
2. report.pdf - your project report in pdf format

Report Checklist

Your report should briefly describe your implemented solution and fully answer the questions for every part of the assignment. Your description should focus on the most "interesting" aspects of your solution, i.e., any non-obvious implementation choices and parameter settings, and what you have found to be especially important for getting good performance. Feel free to include pseudocode or figures if they are needed to clarify your approach. Your report should be self-contained, and it should (ideally) make it possible for us to understand your solution without having to run your source code.

Kindly structure the report as follows:

1. **Title Page:**

List of all team members, course number and section for which each member is registered, date on which the report was written

2. **Section I:**

Algorithms (Search). This section should describe algorithms and data structures used for all four search strategies. Answer questions like: what is a state? What is the frontier? Do you maintain an explored states list? How are repeated states detected and managed?

3. **Section II:**

Algorithms (A* and Greedy BFS). This section should describe the heuristic(s) used for A* and Greedy BFS, for both the single dot and multiple-dot situations Provide proof that the heuristics for A* are admissible.

4. **Section III:**

Results (Basic Pathfinding). For every algorithm in part 1, (DFS, BFS, Greedy, A*), and every one of the mazes, (medium, big, open), give the maze screenshot with the computed path, the solution cost and the number of expanded nodes (12 cases total)

5. **Section IV:**

Results (Search with multiple dots). For part 2, for each of three mazes (tiny, small, medium), give the solution screenshot, solution cost, and number of expanded nodes for your A* algorithm.

6. **Extra Credit:**

If you have done any work which you think should get extra credit, describe it here.

7. **Statement of Contribution:**

Specify which team-member performed which task. You are encouraged to make this a many-to-many mapping, if applicable. e.g., You can say that "Rahul and Jason both implemented the BFS function, their results were compared for debugging and Rahul's code was submitted. Jason and Mark both implemented the DFS function, Mark's code never ran successfully, so Jason's code was submitted. Section I of the report was written by all 3 team members. Section II by Mark and Jason, Section III by Rahul and Jason." ... and so on.

WARNING: You will not get credit for any solutions that you have obtained, but not included in your report! Make only ONE submission per team.

Only attach files that are the required deliverables in blackboard. Your report must be a formatted pdf document. Pictures and example outputs should be incorporated into the document. Exception: items which are very large or unsuitable for inclusion in a pdf document (e.g. videos or animated gifs) may be put on the web and a URL included in your report.

Extra credit:

We reserve the right to give **bonus points** for any advanced exploration or especially challenging or creative solutions that you implement. This includes, but is not restricted to, the extra credit suggestion given above.